

WEEK 1 - A
2024

DATE: 02-12-

AIR QUALITY INDEX ANALYSIS

AIM:

To write a program to Analysis the Air Quality Index.

CODE:

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
from statsmodels.tsa.seasonal import seasonal_decompose

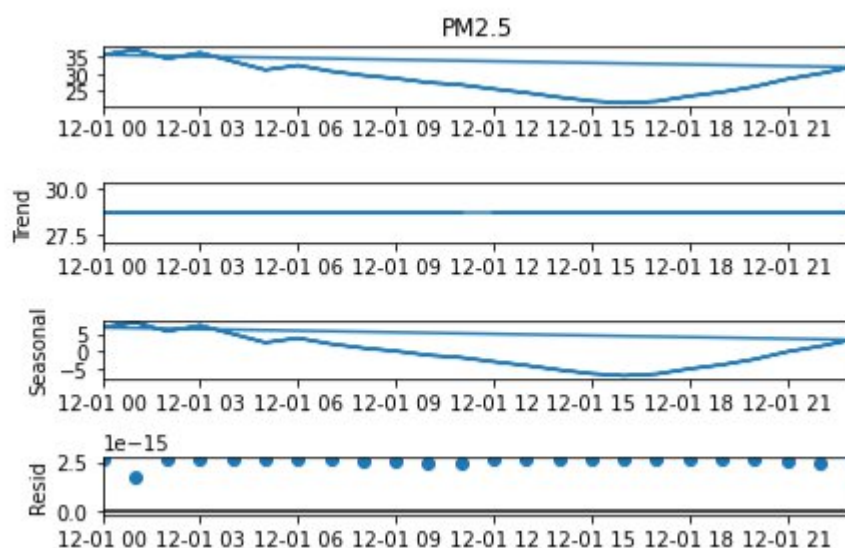
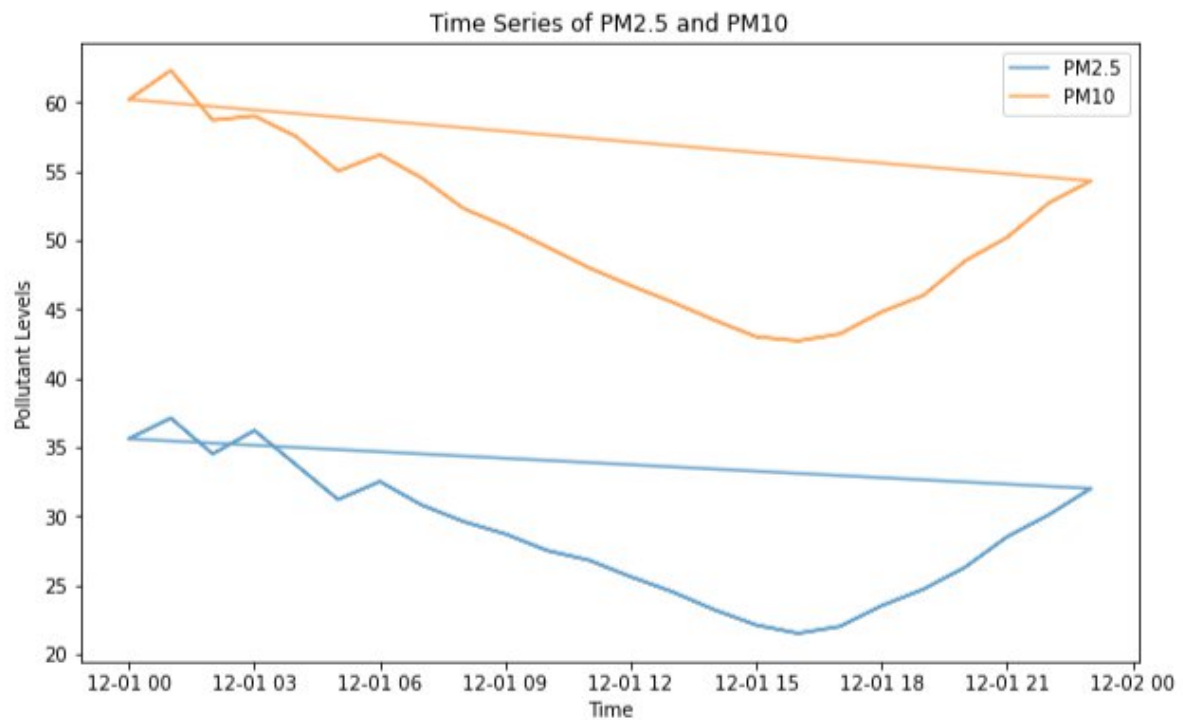
data = {
    "date": ["2024-12-01"] * 24,
    "time": [f"{i:02}:00" for i in range(24)],
    "location": ["CityA"] * 24,
    "PM2.5": [35.6, 37.1, 34.5, 36.2, 33.7, 31.2, 32.5, 30.8, 29.6, 28.7, 27.5, 26.8, 25.6, 24.5, 23.2,
22.1, 21.5, 22.0, 23.5, 24.7, 26.3, 28.5, 30.1, 32.0],
    "PM10": [60.2, 62.3, 58.7, 59.0, 57.5, 55.0, 56.2, 54.5, 52.3, 51.0, 49.5, 48.0, 46.7, 45.5, 44.2,
43.0, 42.7, 43.2, 44.8, 46.0, 48.5, 50.2, 52.7, 54.3],
    "temperature": [22.5, 22.1, 21.9, 21.5,
21.2, 21.0, 21.1, 21.0, 22.0, 23.0, 24.2, 25.1, 26.0, 27.0, 28.0, 28.5, 27.5, 26.5, 25.2,
24.0, 23.0, 22.5, 22.1, 22.0],
    "wind_speed": [3.2, 3.0, 3.1, 3.3, 3.4, 3.5, 3.2, 3.1, 3.0, 2.9, 3.1, 3.2, 3.4, 3.5, 3.6, 3.8, 3.7, 3.5, 3.3,
3.2, 3.0, 2.8, 2.7, 2.5],
}

df = pd.DataFrame(data)
# Convert date and time columns into a single datetime column
df['datetime'] = pd.to_datetime(df['date'] + ' ' + df['time'])
df.set_index('datetime', inplace=True)
# Expand the dataset to simulate two full cycles
df_extended = pd.concat([df, df.copy()], ignore_index=False)
```

```
# Handle missing data (if any)
df_extended.fillna(df_extended.select_dtypes(include=np.number).mean(), inplace=True)

# Plot time-series data
plt.figure(figsize=(10, 6))
plt.plot(df_extended.index, df_extended['PM2.5'], label='PM2.5', alpha=0.7)
plt.plot(df_extended.index, df_extended['PM10'], label='PM10', alpha=0.7)
plt.title('Time Series of PM2.5 and PM10')
plt.xlabel('Time')
plt.ylabel('Pollutant Levels')
plt.legend()
plt.show()

# Decompose time series for seasonality analysis
result = seasonal_decompose(df_extended['PM2.5'], model='additive', period=24) # Period is 24
hours
result.plot()
plt.show()
```

OUTPUT:**RESULT:**

Thus, The program is successfully executed.

WEEK 1 - B
2024

DATE: 02-12-

E-COMMERCE SALES TREND

AIM:

To write a program to E-Commerce Sales Trends Analysis.

CODE:

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
from statsmodels.tsa.seasonal import seasonal_decompose
# Load the dataset into a DataFrame
sales_data = {
    "Product Category": ["Electronics", "Clothing", "Groceries", "Electronics", "Clothing",
                        "Groceries",
                        "Electronics", "Clothing", "Groceries", "Electronics", "Clothing", "Groceries",
                        "Electronics", "Clothing", "Groceries", "Electronics", "Clothing", "Groceries",
                        "Electronics", "Clothing", "Groceries"],
    "Sales Volume": [120, 85, 240, 150, 100, 220, 200, 120, 300, 140, 90, 260, 175, 110, 270, 190,
                    100, 250, 160, 120, 280],
    "Price": [2500, 1500, 500, 2700, 1700, 520, 3000, 1800, 480, 2500, 1600, 450, 2800, 1750, 520,
             2900, 1650, 500, 2650, 1800, 530],
    "Customer Demographics": ["Urban", "Rural", "Urban", "Suburban", "Urban", "Rural",
                             "Urban", "Urban", "Rural", "Suburban", "Rural", "Urban",
                             "Urban", "Suburban", "Rural", "Rural", "Urban", "Urban",
                             "Suburban", "Rural", "Urban"],
    "Time": ["2024-12-01", "2024-12-01", "2024-12-01", "2024-12-02", "2024-12-02", "2024-12-02",
            "2024-12-03", "2024-12-03", "2024-12-03", "2024-12-04", "2024-12-04", "2024-12-04",
            "2024-12-05", "2024-12-05", "2024-12-05", "2024-12-06", "2024-12-06", "2024-12-06",
            "2024-12-07", "2024-12-07", "2024-12-07"],
}
```

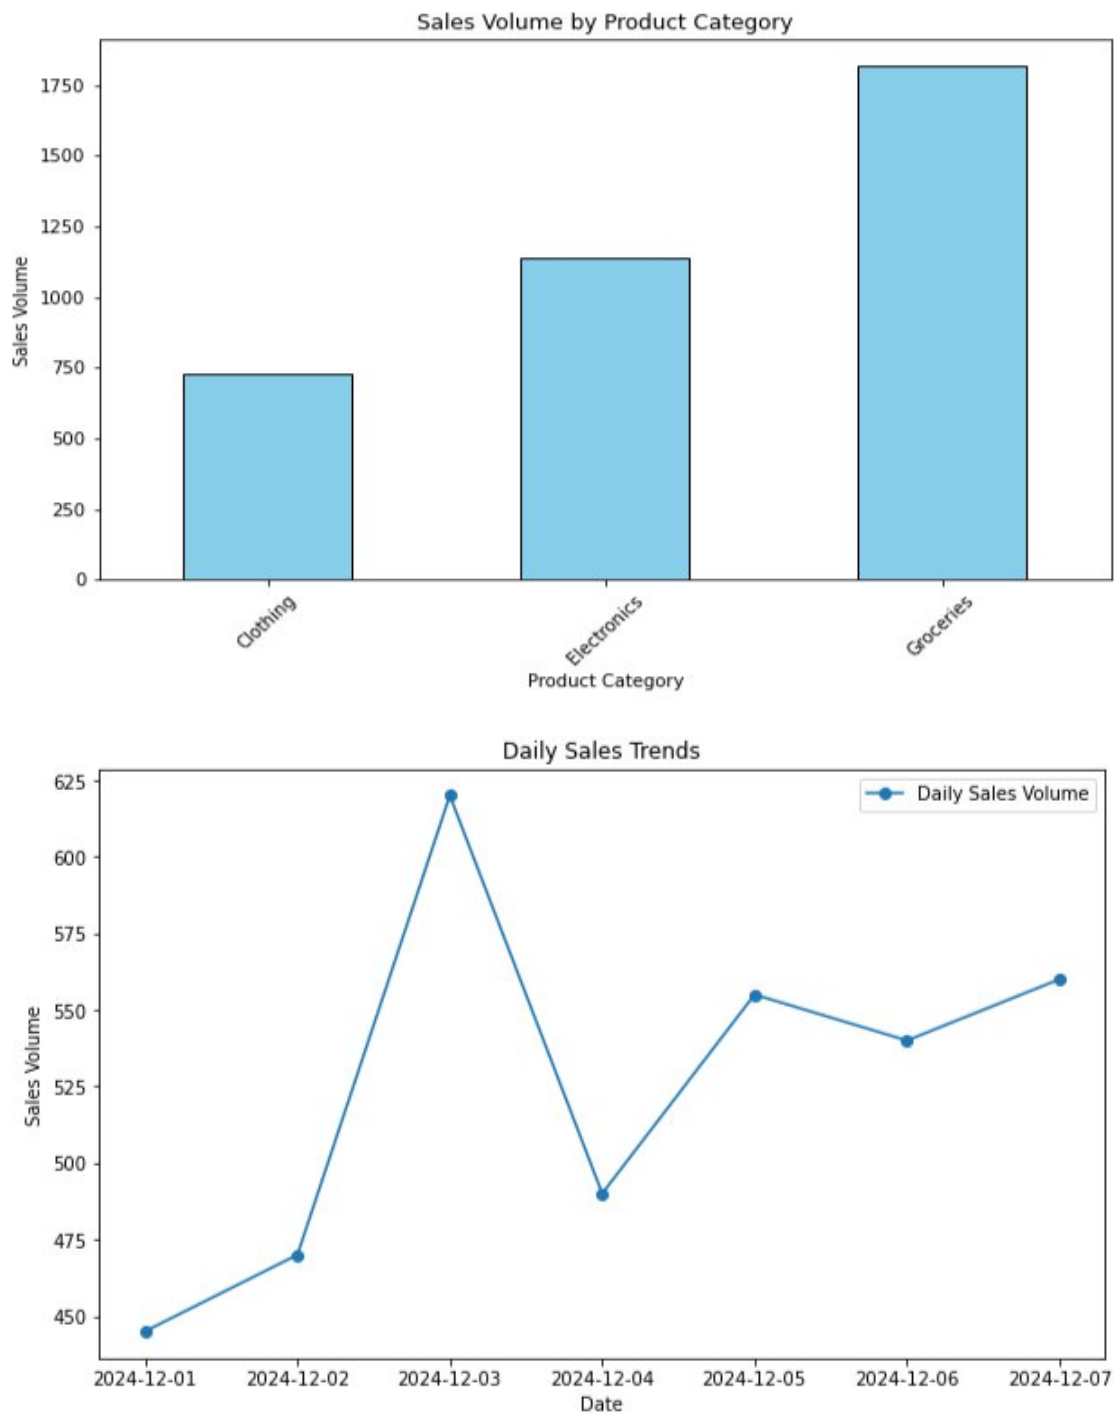
```
sales_df = pd.DataFrame(sales_data)

# Convert 'Time' column to datetime
sales_df['Time'] = pd.to_datetime(sales_df['Time'])

# Group data by product category and aggregate sales volume and price
category_sales = sales_df.groupby('Product Category')[['Sales Volume', 'Price']].sum()

# Plot sales volume by product category
plt.figure(figsize=(10, 6))
category_sales['Sales Volume'].plot(kind='bar', color='skyblue', edgecolor='black')
plt.title('Sales Volume by Product Category')
plt.xlabel('Product Category')
plt.ylabel('Sales Volume')
plt.xticks(rotation=45)
plt.show()

# Group by time and visualize sales trends
time_sales = sales_df.groupby('Time')['Sales Volume'].sum()
plt.figure(figsize=(10, 6))
plt.plot(time_sales.index, time_sales, marker='o', label='Daily Sales Volume')
plt.title('Daily Sales Trends')
plt.xlabel('Date')
plt.ylabel('Sales Volume')
plt.legend()
plt.show()
```

OUTPUT:

RESULT:

Thus, The program is successfully executed.

WEEK 1 - C
2024

DATE: 02-12-

COVID - 19 CASE STUDY

AIM:


To write a program to Analysis the Covid-19 Case Study.

CODE:

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.ensemble import IsolationForest
# Example COVID-19 dataset
data = {
    'Date': pd.date_range(start='2024-01-01', periods=10, freq='D'),
    'Region': ['Region A']*5 + ['Region B']*5,
    'Confirmed Cases': [100, 150, 200, 250, 300, 50, 60, 70, 80, 90],
    'Recovered Cases': [50, 80, 120, 150, 200, 30, 40, 50, 60, 70],
    'Deaths': [5, 7, 10, 12, 15, 2, 3, 4, 5, 6]
}
# Convert dictionary into a pandas DataFrame
df = pd.DataFrame(data)
# Calculate rolling averages (3-day window for example)
df['Rolling Confirmed'] = df.groupby('Region')['Confirmed Cases'].rolling(window=3).mean().reset_index(0, drop=True)
df['Rolling Recovered'] = df.groupby('Region')['Recovered Cases'].rolling(window=3).mean().reset_index(0, drop=True)
df['Rolling Deaths'] = df.groupby('Region')['Deaths'].rolling(window=3).mean().reset_index(0, drop=True)
```



```
# Prepare data for anomaly detection
features = ['Confirmed Cases', 'Recovered Cases', 'Deaths'] # Using raw features for simplicity
df_for_model = df[features]
iso_forest = IsolationForest(contamination=0.1, random_state=42) # Adjust contamination as needed
df['Anomaly'] = iso_forest.fit_predict(df_for_model)
# Map anomaly values to more intuitive labels
df['Anomaly'] = df['Anomaly'].map({1: 'Normal', -1: 'Anomaly'})
# Visualize the results
plt.figure(figsize=(12, 6))
sns.scatterplot(
    data=df,
    x='Date',
    y='Confirmed Cases',
    hue='Anomaly',
    style='Region',
    palette={'Normal': 'blue', 'Anomaly': 'red'},
    s=100
)
plt.title('Anomaly Detection in Confirmed Cases')
plt.xticks(rotation=45)
plt.show()
# Print the DataFrame to verify anomalies
print(df)
```

OUTPUT:

	Date	Region	Confirmed Cases	Recovered Cases	Deaths	\
0	2024-01-01	Region A	100	50	5	
1	2024-01-02	Region A	150	80	7	
2	2024-01-03	Region A	200	120	10	
3	2024-01-04	Region A	250	150	12	
4	2024-01-05	Region A	300	200	15	
5	2024-01-06	Region B	50	30	2	
6	2024-01-07	Region B	60	40	3	
7	2024-01-08	Region B	70	50	4	
8	2024-01-09	Region B	80	60	5	
9	2024-01-10	Region B	90	70	6	

	Rolling Confirmed	Rolling Recovered	Rolling Deaths	Anomaly
0	NaN	NaN	NaN	Normal
1	NaN	NaN	NaN	Normal
2	150.0	83.333333	7.333333	Normal
3	200.0	116.666667	9.666667	Normal
4	250.0	156.666667	12.333333	Anomaly
5	NaN	NaN	NaN	Normal
6	NaN	NaN	NaN	Normal
7	60.0	40.000000	3.000000	Normal
8	70.0	50.000000	4.000000	Normal
9	80.0	60.000000	5.000000	Normal

RESULT:

Thus, The program is successfully executed.

WEEK 2 - A
2024

DATE: 10-12-

SPAM (OR) NOT SPAM

AIM:

To write a program to classify the email as spam or not spam using KNN neighbour algorithm.

CODE:

```
#Import the needed libraries
import pandas as pd
import numpy as np
from sklearn.model_selection import train_test_split
from sklearn.feature_extraction.text import TfidfVectorizer
from sklearn.naive_bayes import MultinomialNB
from sklearn.metrics import accuracy_score, classification_report, confusion_matrix
data=pd.read_csv('dataset.csv')
print(data.head())
data.dropna(inplace=True)
data['text']=data['text'].str.lower()
X=data['text']
Y=data['label']
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)
vectorizer = TfidfVectorizer(stop_words='english', max_features=5000)
X_train_tfidf = vectorizer.fit_transform(X_train)
X_test_tfidf = vectorizer.transform(X_test)
model = MultinomialNB()
model.fit(X_train_tfidf, y_train)
y_pred = model.predict(X_test_tfidf)
print("Accuracy:", accuracy_score(y_test, y_pred))
print("\nClassification Report:\n", classification_report(y_test, y_pred))
```

```
print("\nConfusion Matrix:\n", confusion_matrix(y_test, y_pred))
new_email = ["Congratulations! You've won a $1000 gift card. Click here to claim."]
new_email_tfidf = vectorizer.transform(new_email)
prediction = model.predict(new_email_tfidf)
print("\nPrediction for new email:", "Scam" if prediction[0] == 1 else "Not Scam")
```

OUTPUT:

Cross-Validation Accuracy: 0.87

Test Accuracy: 0.60

Classification Report:

	precision	recall	f1-score	support
0	1.00	0.50	0.67	4
1	0.33	1.00	0.50	1
accuracy			0.60	5
macro avg	0.67	0.75	0.58	5
weighted avg	0.87	0.60	0.63	5

Model saved as 'spam_classifier_model.pkl'

Prediction for test email (0 = Not Spam, 1 = Spam): 1

RESULT:

Thus, The program is successfully executed.

WEEK 2 – B
2024

DATE: 10-12-

PIZZA LIKING PREDICTION

AIM:

To write a program to Pizza Liking Prediction using KNN neighbour algorithm.

CODE:

```
import pandas as pd
import seaborn as sns
from sklearn.model_selection import train_test_split
from sklearn.ensemble import RandomForestClassifier
from sklearn.metrics import accuracy_score, classification_report

data={ "age":[15,25,35,45,55,22,40,28,18,60],
       "likes_chease":[1,1,0,1,0,1,0,1,1,0],
       "Fast_food":[0,0,1,0,1,0,1,0,0,1],
       "Salary":[50000,100000,10000,80000,8000,45000,11000,120000,130000,13000],
       "location":["Urban","Urban","Rural","Urban","Rural","Urban","Rural","Urban","Urban","Rural"],
       "Awareness":["know","know","Unknow","know","Unknow","know","Unknow","know","know","Unknow"],
       "working":["IT","IT","Farmer","Engineer","Driver","Pilot","Painter","IT","Artisr","Player"],
       "Health":["Good","Bad","Bad","Good","Bad","Good","Bad","Good","Good","Bad"],
       "likes_pizza":["yes","yes","no","yes","no","yes","no","yes","yes","no"]}

df=pd.DataFrame(data)
x=df[["age","likes_chease","Fast_Food","Salary","location","working","Health","Awareness"]]
y=df["likes_pizza"]
x_train,x_test,y_train,y_test=train_test_split(x, y, test_size=0.3, random_state=42)
model=RandomForestClassifier(random_state=42)
model.fit(x_train,y_train)
y_pred = model.predict(x_test)
print(df)
print("Accuracy:",accuracy_score(y_test,y_pred))
print("\n Classification Report:\n",classification_report(y_test,y_pred))
example_person={"age":30,"likes_cheese":1,"vegetarian":0,"like_spicy":1}
example_df=pd.DataFrame([example_person])
prediction=model.predict(example_df)
```

```
print(f"\nPrediction: The person {'likes' if prediction[0] == 'yes' else 'does not like'} pizza.")
```

OUTPUT:

```

    age  likes_chease  Fast_food  Salary location Awareness  working
Health \
0   15              1          0   50000   Urban      know      IT
Good
1   25              1          0  100000   Urban      know      IT
Bad
2   35              0          1   10000   Rural     Unknow    Farmer
Bad
3   45              1          0   80000   Urban      know    Engineer
Good
4   55              0          1    8000   Rural     Unknow    Driver
Bad
5   22              1          0   45000   Urban      know      Pilot
Good
6   40              0          1   11000   Rural     Unknow    Painter
Bad
7   28              1          0  120000   Urban      know      IT
Good
8   18              1          0  130000   Urban      know    Artisr
Good
9   60              0          1   13000   Rural     Unknow    Player
Bad
likes_pizza
0          yes
1          yes
2          no
3          yes
4          no
5          yes
6          no
7          yes
8          yes
9          no
Accuracy: 1.0

```

```

Classification Report:
              precision    recall  f1-score   support

    yes               1.00      1.00      1.00         3

   accuracy               1.00      1.00      1.00         3
  macro avg               1.00      1.00      1.00         3
weighted avg               1.00      1.00      1.00         3

```

```
Prediction: The person likes pizza.
```

RESULT:

Thus, The program is successfully executed.

WEEK 2 – C
2024

DATE: 10-12-

MOVIE GENRE PREDICTION

AIM:

To write a program to Analysis the Movie Genre Prediction.

CODE:

```
import pandas as pd

from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler, LabelEncoder
from sklearn.ensemble import RandomForestClassifier
from sklearn.metrics import classification_report

# Load the dataset
df = pd.read_csv("movie.csv")

# Encode categorical features
label_encoder = LabelEncoder()
df['language'] = label_encoder.fit_transform(df['language'])
df['genre'] = label_encoder.fit_transform(df['genre'])
df['director'] = label_encoder.fit_transform(df['director'])

# Remove rare classes with fewer than 2 samples
class_counts = df['genre'].value_counts()
rare_classes = class_counts[class_counts < 2].index
df = df[~df['genre'].isin(rare_classes)]

# Features and target
X = df[['duration', 'language', 'average_rating', 'number_of_reviews', 'year', 'budget', 'revenue']]
y = df['genre']

# Check class distribution
print("Class distribution in the target variable:")
print(df['genre'].value_counts())

# Scale features
```

```

scaler = StandardScaler()
X_scaled = scaler.fit_transform(X)
# Split the data with stratification
X_train, X_test, y_train, y_test = train_test_split(X_scaled, y, test_size=0.3, random_state=42,
stratify=y)
# Train a classifier with class weights to handle imbalance
clf = RandomForestClassifier(random_state=42, class_weight="balanced")
clf.fit(X_train, y_train)
# Predictions
y_pred = clf.predict(X_test)
# Evaluate using classification report with zero_division parameter
print("Classification Report:")
print(classification_report(y_test, y_pred, zero_division=0))

```

OUTPUT:

```

Class distribution in the target variable:
1      3
0      2
Name: genre, dtype: int64
Classification Report:

```

	precision	recall	f1-score	support
0	0.00	0.00	0.00	1
1	0.50	1.00	0.67	1
accuracy			0.50	2
macro avg	0.25	0.50	0.33	2
weighted avg	0.25	0.50	0.33	2

RESULT:

Thus, The program is successfully executed.

WEEK 2 – D
2024

DATE: 10-12-

SPORTS PERFORMANCE ANALYSIS

AIM:

To write a program to Analysis the Sports Performance.

CODE:

```
import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.neighbors import KNeighborsClassifier
from sklearn.metrics import classification_report

# Load the dataset
df = pd.read_csv("sports.csv")

# Features and target
X = df[['accuracy', 'speed', 'stamina', 'age']].copy() # Ensure X is a copy
y = df['performance'].copy() # Ensure y is a copy

# Include outliers
outlier = pd.DataFrame([[200, 15, 150, 30]], columns=X.columns)
X = pd.concat([X, outlier], ignore_index=True) # Add a corresponding target value for the outlier
using pd.concat
y = pd.concat([y, pd.Series(['excellent'])], ignore_index=True)

# Split the data
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3, random_state=42)

# Train k-NN model
knn = KNeighborsClassifier(n_neighbors=3)
knn.fit(X_train, y_train)

# Predictions
y_pred = knn.predict(X_test)
```

```
# Evaluate with zero_division set to 1 to avoid warnings  
print(classification_report(y_test, y_pred, zero_division=1))
```

OUTPUT:

	precision	recall	f1-score	support
Average	0.00	1.00	0.00	0.0
Excellent	1.00	0.00	0.00	1.0
Good	1.00	0.00	0.00	1.0
accuracy			0.00	2.0
macro avg	0.67	0.33	0.00	2.0
weighted avg	1.00	0.00	0.00	2.0

RESULT:

Thus, The program is successfully executed.

WEEK 3 – A
2024

DATE: 19-12-

FUEL AMOUNT PREDICTION

AIM:

To write a program to Analysis the fuel amount prediction using linear regression.

CODE:

```
import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
from sklearn.metrics import mean_squared_error
df=pd.read_csv("Fuel_data.csv")
print("Dataset Preview:")
print(df.head())
X = df[['distance']] # Feature: distance traveled
y = df['fuel'] # Target: fuel consumed
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)
model = LinearRegression()
model.fit(X_train, y_train)
y_pred = model.predict(X_test)
mse = mean_squared_error(y_test, y_pred)
print(f"Mean Squared Error: {mse}")
print("Model Coefficients:")
print(f"Intercept: {model.intercept_}")
print(f"Slope: {model.coef_[0]}")
new_distance = pd.DataFrame({'distance': [150]}) # Create DataFrame for new input
predicted_fuel = model.predict(new_distance)
print(f"Predicted fuel for {new_distance.iloc[0, 0]} km: {predicted_fuel[0]:.2f} liters")
```

OUTPUT:

```
Dataset Preview:
  distance  fuel
0         10   0.8
1         20   1.6
2         30   2.4
3         40   3.2
4         50   4.0
Mean Squared Error: 9.860761315262648e-32
Model Coefficients:
Intercept: 8.881784197001252e-16
Slope: 0.07999999999999999
Predicted fuel for 150 km: 12.00 liters
```

RESULT:

Thus, The program is successfully executed.

WEEK 3 – B
2024

DATE: 19-12-

SALARY PREDICTION

AIM:

To write a program to Analysis the fuel amount prediction using linear regression.

CODE:

```
import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
from sklearn.metrics import mean_squared_error
df = pd.read_csv("salary_data.csv")
df = pd.get_dummies(df, columns=['industry', 'location'], drop_first=True)
X = df[['years_experience', 'qualification']] + [col for col in df.columns if 'industry_' in col or 'location_' in col]
y = df['salary']
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)
model = LinearRegression()
model.fit(X_train, y_train)
y_pred = model.predict(X_test)
print("Dataset Preview:")
print(df.head())
mse = mean_squared_error(y_test, y_pred)
print(f"Mean Squared Error: {mse}")
new_data = pd.DataFrame({ 'years_experience': [5],
    'qualification': [3],
    'industry_IT': [1],
    'industry_Marketing': [0], # Ensure this matches training columns
    'location_CityB': [1],
    'location_CityC': [0],
}, columns=X_train.columns) # Ensure exact column alignment
```

```
predicted_salary = model.predict(new_data)
print(f'Predicted salary for the new profile: {predicted_salary[0]:.2f}')
```

OUTPUT:

Dataset Preview:

Years experience	qualification	salary	industry IT	industry Marketing
0	2	2	50000	1
0				
1	5	3	80000	0
0				
2	3	2	55000	1
0				
3	7	4	90000	0
0				
4	1	1	45000	0
1				

	location_CityB	location_CityC
0	0	0
1	1	0
2	0	0
3	1	0
4	0	1

Mean Squared Error: 4674945.215485762

Predicted salary for the new profile: 72162.16

RESULT:

Thus, The program is successfully executed.

WEEK 4 – A
2025

DATE: 24-01-

HOUSE PRICE PREDICTION

AIM:

To write a program to Analysis the House Price Prediction using L-R regulation.

CODE:

```
import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.linear_model import Ridge
from sklearn.metrics import mean_squared_error
data = pd.read_csv("house_prices.csv")
X = data[['size', 'bedrooms', 'age']]
y = data['price']
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)
model = Ridge(alpha=1.0) # Alpha controls the regularization strength
model.fit(X_train, y_train)
y_pred = model.predict(X_test)
print("Mean Squared Error:", mean_squared_error(y_test,y_pred))
```

OUTPUT:

Mean Squared Error: 3325492.320725987

RESULT:

Thus, The program is successfully executed.

WEEK 4 – B
2025

DATE: 24-01-

CROP YIELD PREDICTION

AIM:

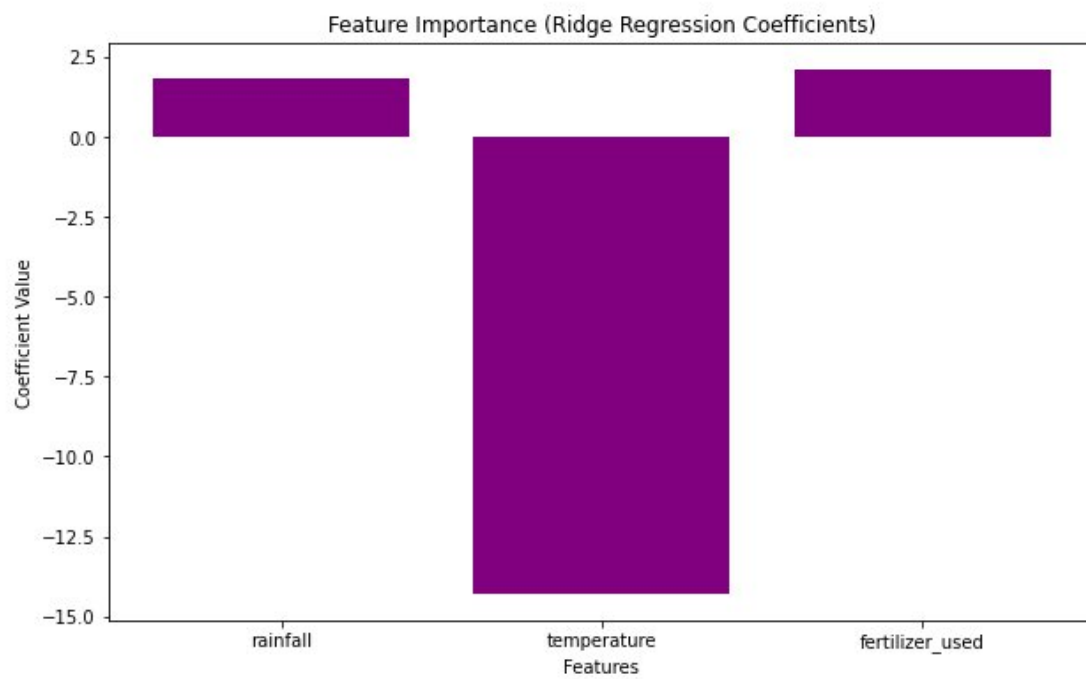
To write a program to Analysis the Crop Yield Prediction using L-R regulation.

CODE:

```
import pandas as pd
import matplotlib.pyplot as plt
from sklearn.model_selection import train_test_split
from sklearn.linear_model import Ridge
from sklearn.metrics import mean_squared_error
data = pd.read_csv("crop_yield.csv")
X = data[['rainfall', 'temperature', 'fertilizer_used']]
y = data['yield']
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)
model = Ridge(alpha=1.0)
model.fit(X_train, y_train)
y_pred = model.predict(X_test)
mse = mean_squared_error(y_test, y_pred)
print("Mean Squared Error:", mse)
coefficients = model.coef_
plt.figure(figsize=(10, 6))
plt.bar(X.columns, coefficients, color='purple')
plt.title('Feature Importance (Ridge Regression Coefficients)')
plt.xlabel('Features')
plt.ylabel('Coefficient Value')
plt.show()
```

OUTPUT:

Mean Squared Error: 2498.780533491533

**RESULT:**

Thus, The program is successfully executed.

WEEK 4 – C
2025

DATE: 24-01-

ENERGY EFFICIENCY PREDICTION

AIM:

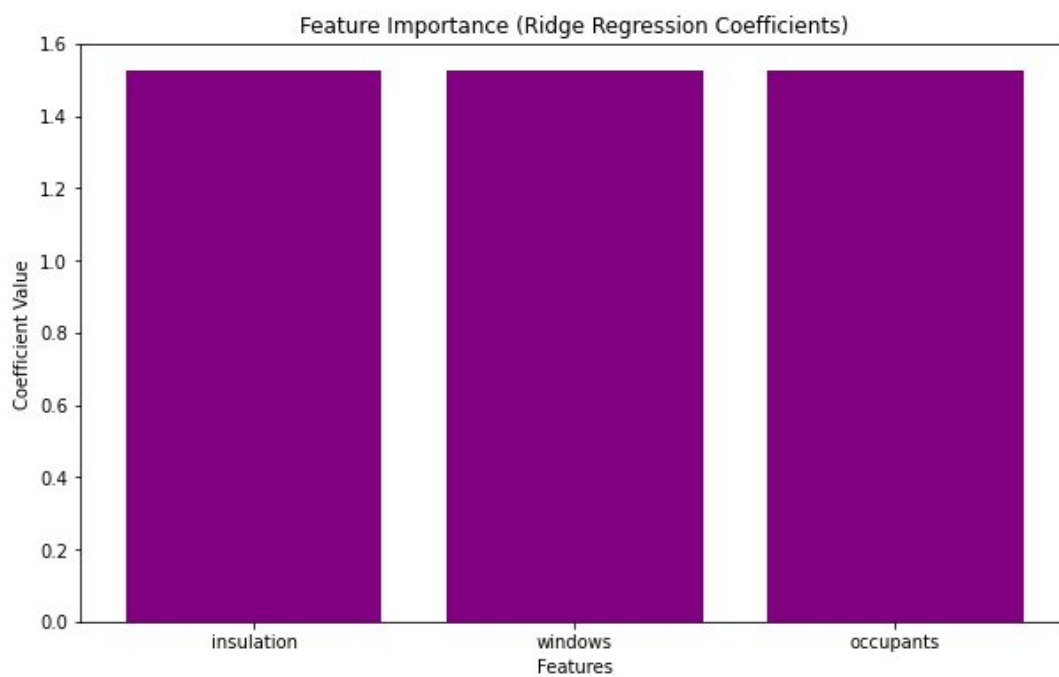
To write a program to Analysis the Energy Efficiency Prediction using L-R regulation.

CODE:

```
import pandas as pd
import matplotlib.pyplot as plt
from sklearn.model_selection import train_test_split
from sklearn.linear_model import Ridge
from sklearn.metrics import mean_squared_error
data = pd.read_csv("energy_efficiency.csv")
X = data[['insulation', 'windows', 'occupants']]
y = data['efficiency']
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)
model = Ridge(alpha=1.0)
model.fit(X_train, y_train)
y_pred = model.predict(X_test)
print("Mean Squared Error:", mean_squared_error(y_test, y_pred))
coefficients = model.coef_
plt.figure(figsize=(10, 6))
plt.bar(X.columns, coefficients, color='purple')
plt.title('Feature Importance (Ridge Regression Coefficients)')
plt.xlabel('Features')
plt.ylabel('Coefficient Value')
plt.show()
```

OUTPUT:

Mean Squared Error: 0.03251814028486949

**RESULT:**

Thus, The program is successfully executed.

WEEK 5 – A
2025

DATE: 24-01-

DIABETES CLASSIFICATION

AIM:

To write a program to Diabetes Classification using L-R regulation.

CODE:

```
import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import accuracy_score, classification_report
data = pd.read_csv("diabetes.csv")
X = data[['age', 'bmi', 'blood_pressure', 'glucose']]
y = data['diabetes']
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)
model = LogisticRegression()
model.fit(X_train, y_train)
y_pred = model.predict(X_test)
print("Accuracy:", accuracy_score(y_test, y_pred))
print("Classification Report:\n", classification_report(y_test, y_pred))
```

OUTPUT:

```
Accuracy: 1.0
Classification Report:
              precision    recall  f1-score   support

0               1.00      1.00      1.00         1

   accuracy                1.00                1
  macro avg               1.00                1
weighted avg               1.00                1
```

RESULT:

Thus, The program is successfully executed.

WEEK 5 – B
2025

DATE: 24-01-

CREDIT CARD DEFAULT PREDICTION

AIM:

To write a program to Analysis the Credit Card Default Prediction using L-R regulation.

CODE:

```
import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import accuracy_score, classification_report
data = pd.read_csv("credit_card_default.csv")
X = data[['income', 'credit_score', 'age', 'loan_amount']]
y = data['default']
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)
model = LogisticRegression()
model.fit(X_train, y_train)
y_pred = model.predict(X_test)
print("Accuracy:", accuracy_score(y_test, y_pred))
print("Classification Report:\n", classification_report(y_test,y_pred))
```

OUTPUT:

```
Accuracy: 0.0
Classification Report:
              precision    recall  f1-score   support

     0              0.00      0.00      0.00         0.0
     1              0.00      0.00      0.00         1.0

 accuracy                0.00         0.00         1.0
 macro avg              0.00      0.00      0.00         1.0
 weighted avg          0.00      0.00      0.00         1.0
```

RESULT:

Thus, The program is successfully executed.

WEEK 5 – C
2025

DATE: 24-01-

HEART DISEASE CLASSIFICATION

AIM:

To write a program to Analysis the Heart Disease Classification using L-R regulation.

CODE:

```
import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import accuracy_score, classification_report
data = pd.read_csv("heart_disease.csv")
X = data[['age', 'cholesterol', 'blood_pressure', 'exercise']]
y = data['heart_disease']
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)
model = LogisticRegression()
model.fit(X_train, y_train)
y_pred = model.predict(X_test)
print("Accuracy:", accuracy_score(y_test, y_pred))
print("Classification Report:\n", classification_report(y_test,y_pred))
```

OUTPUT:

```
Accuracy: 1.0
Classification Report:
              precision    recall  f1-score   support

0               1.00      1.00      1.00         1

   accuracy                1.00         1
  macro avg               1.00      1.00      1.00         1
 weighted avg               1.00      1.00      1.00         1
```

RESULT:

Applied Machine Learning Lab (P24DS2P6)

Thus, The program is successfully executed.

Week-6A**Date: 31/01/2025****PREDICTIVE ANALYTICS FOR HOSPITALS
DISEASE OUTBREAK PREDICATION****AIM:**

To write a program to Disease outbreak predication

PROGRAM CODE:

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
from sklearn.linear_model import LinearRegression
from sklearn.preprocessing import PolynomialFeatures
from sklearn.pipeline import make_pipeline
from sklearn.model_selection import train_test_split
from sklearn.metrics import mean_absolute_error, r2_score

# Load data
df = pd.read_csv('disease_outbreak_data.csv')
df['date'] = pd.to_datetime(df['date'])
df['day_number'] = (df['date'] - df['date'].min()).dt.days

X = df[['day_number']]
y = df['number_of_cases']

# Split data
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)

# Use Polynomial Regression (degree=2) for better prediction
model = make_pipeline(PolynomialFeatures(degree=2), LinearRegression())
model.fit(X_train, y_train)

# Predictions
y_pred = model.predict(X_test)

# Evaluate the model
mae = mean_absolute_error(y_test, y_pred)
r2 = r2_score(y_test, y_pred)
```

```
print(f'MAE: {mae:.2f}, R²: {r2:.2f}')

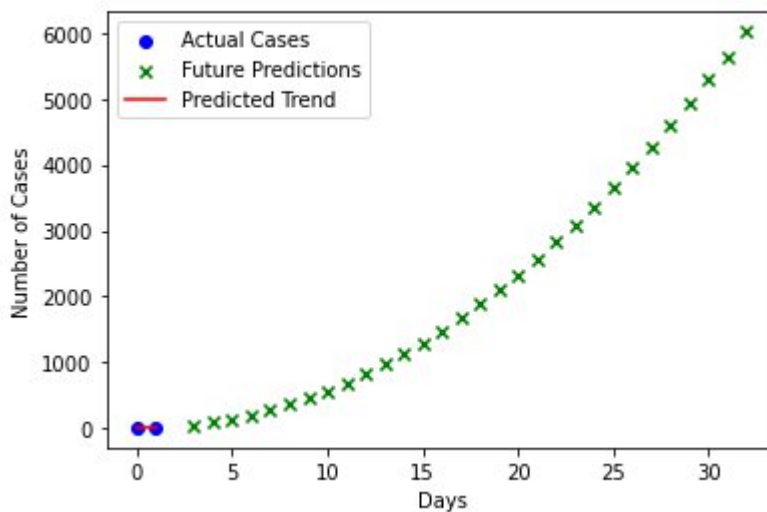
# Forecast future cases (next 30 days)
future_days = pd.DataFrame(np.arange(X.max().values[0] + 1, X.max().values[0] + 31),
                           columns=['day_number'])
future_predictions = model.predict(future_days)

# Sort X_test for correct plotting
X_test_sorted, y_pred_sorted = zip(*sorted(zip(X_test.values.flatten(), y_pred)))

# Plot results
plt.scatter(X_test, y_test, color='blue', label='Actual Cases')
plt.scatter(future_days, future_predictions, color='green', marker='x', label='Future Predictions')
plt.plot(X_test_sorted, y_pred_sorted, color='red', label='Predicted Trend')
plt.xlabel("Days")
plt.ylabel("Number of Cases")
plt.legend()
plt.show()
```

OUTPUT:

MAE: 6.50, R²: -6.12

**RESULT:**

Thus the program executed successfully.

Week-6B**Date: 31/01/2025****BED OCCUPANCY DATA****AIM:**

To write a program to predictive analytics bed occupancy data

PROGRAM CODE:

```
import pandas as pd
import matplotlib.pyplot as plt
from sklearn.ensemble import RandomForestRegressor
from sklearn.model_selection import train_test_split

# Load data
data = pd.read_csv('bed_occupancy_data.csv')

# Ensure 'day' column is numerical
X = data[['day']].values.ravel() # Convert to 1D array
y = data['bed_occupancy']

# Split data
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)


# Reshape X for model input
X_train = X_train.reshape(-1, 1)
X_test = X_test.reshape(-1, 1)

# Model training
rf_model = RandomForestRegressor()
rf_model.fit(X_train, y_train)

# Prediction
predicted_occupancy = rf_model.predict(X_test)

# Plot
plt.scatter(X_test, y_test, color='green', label='Actual')
plt.scatter(X_test, predicted_occupancy, color='orange', label='Predicted') # Use scatter instead
of plot
```

```
plt.xlabel('Day')  
plt.ylabel('Bed Occupancy')  
plt.legend()  
plt.show()
```

OUTPUT:**RESULT:**

Thus the program executed successfully.

Week-6C**Date: 31/01/2025****MEDICATION EFFECTIVENESS DATA****AIM:**

To write a program to predictive analytics medication effectiveness data

PROGRAM CODE:

```
from sklearn.preprocessing import LabelEncoder
from sklearn.tree import DecisionTreeClassifier
# Load data
data = pd.read_csv('medication_effectiveness_data.csv')
# Encode categorical variables
le = LabelEncoder()
data['medication'] = le.fit_transform(data['medication'])
X = data[['medication', 'duration']]
y = data['effectiveness_score'] > 7 # Binary classification: Effective (True/False)
# Split data
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)
# Train classifier
clf = DecisionTreeClassifier()
clf.fit(X_train, y_train)
# Predictions
y_pred = clf.predict(X_test)
print("Predictions (Effective/Not Effective):", y_pred)
```

OUTPUT:

```
Predictions (Effective/Not Effective): [ True]
```

RESULT:

Thus the program executed successfully.

Aim:

To classify loan applications as approved or rejected based on applicant features using Support Vector Machine (SVM).

Program Code:

```
import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler,
LabelEncoder from sklearn.svm import SVC
from sklearn.metrics import accuracy_score,
classification_report # Load the dataset
dataset =
pd.read_csv('loan_approval_data.csv
') # Encode categorical features
manually label_encoder =
LabelEncoder()
dataset['Employment_Status'] =
label_encoder.fit_transform(dataset['Employment_Status']) #
Preprocess the data
X =
dataset.drop('Loan_Status',
axis=1) y =
dataset['Loan_Status']
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3,
random_state=42) scaler = StandardScaler()
X_train =
scaler.fit_transform(X_train)
X_test =
scaler.transform(X_test)
# Train the SVM model
model =
SVC(kernel='linear')
model.fit(X_train, y_train)
# Predict and evaluate
AppliedMachineLearnin
```

```
y_pred = model.predict(X_test)
print("Accuracy:", accuracy_score(y_test, y_pred))
print("Classification Report:\n",
classification_report(y_test, y_pred))
```


Output:

Accuracy: 1.0

Classification Report:

	precision	recall	f1-score	support
0	1.00	1.00	1.00	2
1	1.00	1.00	1.00	4
accuracy		1.00		6
macro avg	1.00	1.00	1.00	6
weighted avg	1.00	1.00	1.00	6

Result:

Thus, the output was executed successfully.

Aim:

To assess the creditworthiness of individuals by classifying them as creditworthy or not using SVM.

Program Code:

```
import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler,
LabelEncoder from sklearn.svm import SVC
from sklearn.metrics import accuracy_score,
classification_report # Load the dataset
dataset =
pd.read_csv('credit_worthiness_data.csv
') # Encode categorical features
manually label_encoder =
LabelEncoder()
dataset['Employment_Status'] =
label_encoder.fit_transform(dataset['Employment_Status']) #
Preprocess the data
X = dataset.drop('Credit_Worthy',
axis=1) y =
dataset['Credit_Worthy']
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3,
random_state=42) scaler = StandardScaler()
X_train =
scaler.fit_transform(X_train)
X_test =
scaler.transform(X_test)
# Train the SVM model
model =
SVC(kernel='linear')
model.fit(X_train, y_train)
# Predict and evaluate
AppliedMachineLearnin
```

```
y_pred = model.predict(X_test)
print("Accuracy:", accuracy_score(y_test, y_pred))
print("Classification Report:\n",
classification_report(y_test, y_pred))
```

Output:

Accuracy: 1.0

Classification Report:

	precision	recall	f1-score	support
0	1.00	1.00	1.00	4
1	1.00	1.00	1.00	3
accuracy		1.00		7
macro avg	1.00	1.00	1.00	7
weighted avg	1.00	1.00	1.00	7

Result:

Thus, the output was executed successfully.

Aim:

To classify news articles as real or fake using SVM based on textual data features.

Program Code:

```
import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.feature_extraction.text import
TfidfVectorizer from sklearn.svm import SVC
from sklearn.metrics import accuracy_score,
classification_report # Load the dataset
dataset =
pd.read_csv('fake_news_data.csv')
# Preprocess the data
X =
dataset['text']
y =
dataset['label']
]
vectorizer =
TfidfVectorizer() X =
vectorizer.fit_transform(
X)
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3,
random_state=42) # Train the SVM model
model = SVC(kernel='linear')
model.fit(X_train, y_train)
# Predict and evaluate
y_pred = model.predict(X_test)
print("Accuracy:", accuracy_score(y_test, y_pred))
print("Classification Report:\n",
classification_report(y_test, y_pred))
```

Output:

Accuracy: 0.6666666666666666

Classification Report:

	precision	recall	f1-score	support
0	0.75	0.75	0.75	4
1	0.50	0.50	0.50	2
accuracy		0.67		6
macro avg	0.62	0.62	0.62	6
weighted avg	0.67	0.67	0.67	6

Result:

Thus, the output executed successfully.
AppliedMachineLearnin

Aim:

To classify animals into various species or categories based on their characteristics using a Decision Tree.

Program Code:

```
import pandas as pd
from sklearn.model_selection import
train_test_split from sklearn.tree import
DecisionTreeClassifier from
sklearn.preprocessing import
LabelEncoder
from sklearn.metrics import accuracy_score,
classification_report # Load dataset
df =
pd.read_csv("animal_classification.c
sv") # Drop non-essential columns
df = df.drop(columns=["name"])
# Encode categorical columns
manually le_diet = LabelEncoder()
df["diet"] =
le_diet.fit_transform(df["diet"])
le_class = LabelEncoder()
df["class"] =
le_class.fit_transform(df["class"]) #
Split features and target
X =
df.drop(columns=["class"
]) y = df["class"]
# Ensure balanced split using stratification
X_train, X_test, y_train, y_test = train_test_split(X, y,
test_size=0.2,random_state=42) # Train Decision Tree model
clf = DecisionTreeClassifier()
clf.fit(X_train, y_train)
```

```
# Predictions
```

```
y_pred =
```

```
clf.predict(X_test) #
```

```
Results
```

```
print("Accuracy:", accuracy_score(y_test, y_pred))
```

```
print("Classification Report:\n", classification_report(y_test, y_pred,  
zero_division=1))
```


Output:

Accuracy: 0.5

Classification Report:

	precision	recall	f1-score	support
0	0.00	1.00	0.00	0
1	1.00	1.00	1.00	1
3	1.00	0.00	0.00	1
accuracy		0.50		2
macro avg	0.67	0.67	0.33	2
weighted avg	1.00	0.50	0.50	2

Result:

Thus, the output was executed successfully.

Aim:

To classify different plant diseases based on plant leaf characteristics using a Decision Tree.

Program Code:

```
import pandas as pd
from sklearn.model_selection import
train_test_split from sklearn.tree import
DecisionTreeClassifier from
sklearn.preprocessing import
LabelEncoder
from sklearn.metrics import accuracy_score,
classification_report # Load dataset
df = pd.read_csv("plant_disease_classification.csv") # Replace
with actual dataset # Strip spaces from column names (to avoid
hidden errors)
df.columns =
df.columns.str.strip() #
Identify categorical columns
categorical_columns = df.select_dtypes(include=["object"]).columns # Select only
non-numeric columns
# Apply Label Encoding to categorical columns
df[categorical_columns] =
df[categorical_columns].apply(LabelEncoder().fit_transform) # Split
features and target
X = df.drop(columns=["disease"]) #
Feature columns y = df["disease"] #
Target column
# Train-test split
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
random_state=42) # Train Decision Tree model
clf = DecisionTreeClassifier()
clf.fit(X_train, y_train)
# Predictions
AppliedMachineLearnin
```

```
y_pred =  
clf.predict(X_test) #  
Results  
print("Accuracy:", accuracy_score(y_test, y_pred))  
print("Classification Report:\n",  
classification_report(y_test, y_pred))
```

Output:

Accuracy: 1.0

Classification Report:

	precision	recall	f1-score	support
--	-----------	--------	----------	---------

1	1.00	1.00	1.00	2
---	------	------	------	---

accuracy			1.00	2
----------	--	--	------	---

macro avg	1.00	1.00	1.00	2
-----------	------	------	------	---

weighted avg	1.00	1.00	1.00	2
--------------	------	------	------	---

Result:

Thus, the output was executed successfully.

AppliedMachineLearnin

Aim:

To classify types of vehicles (e.g., car, truck, bike) based on their features using a Decision Tree.

Program Code:

```
import pandas as pd
from sklearn.model_selection import
train_test_split from sklearn.tree import
DecisionTreeClassifier from
sklearn.preprocessing import
LabelEncoder
from sklearn.metrics import accuracy_score,
classification_report # Load dataset
df = pd.read_csv("vehicle_classification.csv") # Replace with
actual dataset # Strip spaces from column names
df.columns = df.columns.str.strip()
# Apply Label Encoding to all categorical columns
label_encoder = LabelEncoder()
df[df.select_dtypes(include=['object']).columns] =
df.select_dtypes(include=['object']).apply(label_encoder.fit_transform)
# Split features and target
X = df.drop(columns=["vehicle_type"]) #
Feature columns y = df["vehicle_type"] #
Target column
# Train-test split
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
random_state=42) # Train Decision Tree model
clf = DecisionTreeClassifier()
clf.fit(X_train, y_train)
# Predictions
y_pred =
clf.predict(X_test) #
Results
print("Accuracy:", accuracy_score(y_test, y_pred))
```

```
print("Classification Report:\n",  
classification_report(y_test, y_pred))
```

Output:

Accuracy: 1.0

Classification Report:

	precision	recall	f1-score	support
1	1.00	1.00	1.00	1
3	1.00	1.00	1.00	1
accuracy		1.00		2
		0		
macro avg	1.00	1.00	1.00	2
		0		
weighted avg	1.00	1.00	1.00	2
		0		

Result:

Thus, the output was executed successfully.

AppliedMachineLearnin

Aim:

To predict whether an employee will hop to another job based on their work history and personal details using Random Forest.

Program Code:

```
import pandas as pd
from sklearn.model_selection import
train_test_split from sklearn.ensemble
import RandomForestClassifier from
sklearn.preprocessing import
StandardScaler
from sklearn.metrics import classification_report, confusion_matrix,
accuracy_score # Load dataset
df = pd.read_csv("employee_hopping.csv")
# Convert categorical features to numeric using
one-hot encoding df = pd.get_dummies(df,
drop_first=True)
# Handle missing values by filling with the mean for
numeric columns
df.fillna(df.select_dtypes(include=['number']).mean(),
inplace=True) # Split data into features and target
X = df.drop(columns=["Hopped"])
y = df["Hopped"] # Target variable indicating if an
employee hopped jobs # Split into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3,
random_state=42) # Standardize the data
scaler = StandardScaler()
X_train_scaled =
scaler.fit_transform(X_train)
X_test_scaled =
scaler.transform(X_test)
# Train Random Forest model
model = RandomForestClassifier(n_estimators=100, random_state=42)
```



```
model.fit(X_train_scaled, y_train)
```

```
y_pred =
```

```
model.predict(X_test_scaled) #
```

Output results

```
print(f"Accuracy: {accuracy_score(y_test, y_pred)}")
```

```
print("Confusion Matrix:\n", confusion_matrix(y_test,  
y_pred))
```

```
print("Classification Report:\n", classification_report(y_test, y_pred))
```

Output:

Accuracy: 0.5

Confusion Matrix:

```
[[0 1]
```

```
[0 1]]
```

Classification Report:

	precision	recall	f1-score	support
--	-----------	--------	----------	---------

0	0.00	0.00	0.00	1
---	------	------	------	---

1	0.50	1.00	0.67	1
---	------	------	------	---

accuracy		0.5		2
		0		

macro avg	0.25	0.50	0.3	2
			3	

weighted avg	0.25	0.50	0.3	2
			3	

Result:

Thus, the output was executed successfully.

AppliedMachineLearnin

Aim:

To predict the eligibility of employees for promotion based on their work performance and other factors using Random Forest.

Program Code:

```
import pandas as pd
from sklearn.model_selection import
train_test_split from sklearn.ensemble
import RandomForestClassifier from
sklearn.preprocessing import
StandardScaler
from sklearn.metrics import classification_report, confusion_matrix,
accuracy_score # Load dataset
df = pd.read_csv("promotion_eligibility.csv")
# Convert categorical features to numeric using
one-hot encoding df = pd.get_dummies(df,
drop_first=True)
# Handle missing values by filling with the mean for
numeric columns
df.fillna(df.select_dtypes(include=['number']).mean(),
inplace=True) # Split data into features and target
X = df.drop(columns=["Promotion_Eligible"])
y = df["Promotion_Eligible"] # Target variable indicating
eligibility for promotion # Split into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3,
random_state=42) # Standardize the data
scaler = StandardScaler()
X_train_scaled =
scaler.fit_transform(X_train)
X_test_scaled =
scaler.transform(X_test)
# Train Random Forest model
model = RandomForestClassifier(n_estimators=100, random_state=42)
```

```
model.fit(X_train_scaled, y_train)
```

```
y_pred =
```

```
model.predict(X_test_scaled) #
```

Output results

```
print(f"Accuracy: {accuracy_score(y_test, y_pred)}")
```

```
print("Confusion Matrix:\n", confusion_matrix(y_test,  
y_pred))
```

```
print("Classification Report:\n", classification_report(y_test, y_pred))
```

Output:

Accuracy: 1.0

Confusion Matrix:

```
[[1 0]
```

```
[0 2]]
```

Classification Report:

	precision	recall	f1-score	support
--	-----------	--------	----------	---------

0	1.00	1.00	1.00	1
---	------	------	------	---

1	1.00	1.00	1.00	2
---	------	------	------	---

accuracy			1.0	3
			0	

macro avg	1.00	1.00	1.0	3
			0	

weighted avg	1.00	1.00	1.0	3
			0	

Result:

Thus, the output was executed successfully.

Aim:

To detect fraudulent transactions in banking based on transactional features using Random Forest.

Program Code:

```
import pandas as pd
from sklearn.model_selection import
train_test_split from sklearn.ensemble
import RandomForestClassifier from
sklearn.preprocessing import
StandardScaler
from sklearn.metrics import classification_report, confusion_matrix,
accuracy_score # Load dataset
df = pd.read_csv("bank_fraud.csv")
# Convert categorical features to numeric using
one-hot encoding df = pd.get_dummies(df,
drop_first=True)
# Handle missing values by filling with the mean for
numeric columns
df.fillna(df.select_dtypes(include=['number']).mean(),
inplace=True) # Split data into features and target
X = df.drop(columns=["Fraud"])
y = df["Fraud"] # Target variable indicating if a
transaction is fraudulent # Split into training and testing
sets
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3,
random_state=42) # Standardize the data
scaler = StandardScaler()
X_train_scaled =
scaler.fit_transform(X_train)
X_test_scaled =
scaler.transform(X_test)
# Train Random Forest model
AppliedMachineLearnin
```

```
model = RandomForestClassifier(n_estimators=100, random_state=42)
model.fit(X_train_scaled, y_train)
y_pred =
model.predict(X_test_scaled) #
Output results
print(f"Accuracy: {accuracy_score(y_test, y_pred)}")
print("Confusion Matrix:\n", confusion_matrix(y_test,
y_pred))
```

```
print("Classification Report:\n", classification_report(y_test, y_pred))
```

Output:

Accuracy: 0.6666666666666666

Confusion Matrix:

```
[[1 0]
```

```
[1 1]]
```

Classification Report:

	precision	recall	f1-score	support
--	-----------	--------	----------	---------

0	0.50	1.00	0.67	1
---	------	------	------	---

1	1.00	0.50	0.67	2
---	------	------	------	---

accuracy			0.6 7	3
----------	--	--	----------	---

macro avg	0.75	0.75	0.6 7	3
-----------	------	------	----------	---

weighted avg	0.83	0.67	0.6 7	3
-----------------	------	------	----------	---

Result:

Thus, the output was executed successfully.

Aim:

To predict different types of physical activities performed by patients using Gradient Boosting.

Program Code:

```
import pandas as
pd import numpy as
np import xgboost
as xgb
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler,
LabelEncoder from sklearn.metrics import
accuracy_score
# Load dataset
df = pd.read_csv("mhealth.csv") # Update with
actual dataset path # Encode categorical labels
label_encoder = LabelEncoder()
df["Activity"] =
label_encoder.fit_transform(df["Activity"]) #
Define features and labels
X =
df.drop(columns=["Activity"
]) y = df["Activity"]
# Train-test split
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
random_state=42) # Standardization
scaler = StandardScaler()
X_train =
scaler.fit_transform(X_train)
X_test =
scaler.transform(X_test)
# Train XGBoost model
model = xgb.XGBClassifier(n_estimators=100, learning_rate=0.1,
```

```
random_state=42) model.fit(X_train, y_train)
# Predictions
y_pred =
model.predict(X_test) #
Accuracy
accuracy = accuracy_score(y_test,
y_pred) print(f"Model Accuracy:
{accuracy:.2f}")
```

```
# Decode labels back to original values
y_pred_original =
label_encoder.inverse_transform(y_pred)
print("Predicted Activities:", y_pred_original)
```

Output:

Model Accuracy: 0.00

Predicted Activities:

['Walking']

Result:

AppliedMachineLearnin

Thus, the output was executed successfully.

Aim:

To predict the progression of diseases in patients based on clinical data using Gradient Boosting.

Program Code:

```
import pandas as
pd import numpy as
np import xgboost
as xgb
from sklearn.model_selection import
train_test_split from sklearn.preprocessing
import StandardScaler
from sklearn.metrics import
mean_absolute_error, r2_score # Load dataset
df = pd.read_csv("disease_progression.csv") # Update with
actual dataset path # Handle missing values (if any)
df.fillna(df.mean(), inplace=True)
# Define features and target variable
X =
df.drop(columns=["Diabetes_Progressi
on"]) y = df["Diabetes_Progression"]
# Train-test split
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
random_state=42) # Standardization
scaler = StandardScaler()
X_train =
scaler.fit_transform(X_train)
X_test =
scaler.transform(X_test)
# Train XGBoost Regressor
model = xgb.XGBRegressor(n_estimators=200, learning_rate=0.05,
max_depth=5, random_state=42) model.fit(X_train, y_train)
# Predictions
```

```
y_pred =  
model.predict(X_test) #  
Model Evaluation  
mae =  
mean_absolute_error(y_test,  
y_pred) r2 = r2_score(y_test,  
y_pred)
```

1stM.Sc.DataScie

```
print(f"Mean Absolute Error:  
{mae:.2f}") print(f"R2 Score:  
{r2:.2f}")
```

Output:

Mean Absolute Error:
44.98 R² Score: nan

Result:

Thus, the output was executed successfully.

AppliedMachineLearnin

Aim:

To predict whether a health insurance claim will be approved or rejected using Gradient Boosting.

Program Code:

```
import pandas as
pd import numpy as
np
from sklearn.model_selection import
train_test_split from sklearn.ensemble
import AdaBoostClassifier from sklearn.tree
import DecisionTreeClassifier from
sklearn.metrics import accuracy_score
# Load dataset
df = pd.read_csv("health_insurance_claims.csv") # Update with
actual dataset path # Define features and labels
X = df.drop(columns=["Claim_Status"]) # 'Claim_Status' should be
the target variable y = df["Claim_Status"]
# Convert categorical features
if needed X =
pd.get_dummies(X)
# Train-test split
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
random_state=42) # Train AdaBoost model
model = AdaBoostClassifier(estimator=DecisionTreeClassifier(max_depth=2),
n_estimators=100, learning_rate=0.1, random_state=42)
model.fit(X_train,
y_train) # Predictions
on test set
y_pred =
model.predict(X_test) #
Accuracy
accuracy = accuracy_score(y_test,
y_pred) print(f"Model Accuracy:
AppliedMachineLearnin
```



```
{accuracy:.2f}") # Prediction for  
new claim  
new_claim =  
    pd.DataFrame([{"Age  
": 45,  
"BMI": 28.5,
```

```
"Smoker": "Yes",  
"Claim_Amount":  
5000,  
"Hospital_Visits": 2  
}]) # Modify based on actual dataset features  
new_claim = pd.get_dummies(new_claim) # Convert  
categorical features # Align new data with training columns  
new_claim = new_claim.reindex(columns=X.columns,  
fill_value=0) # Predict claim status  
predicted_status =  
model.predict(new_claim)[0] print(f"Predicted  
Claim Status: {predicted_status}") Output:  
Model Accuracy: 1.00  
Predicted Claim Status: Approved
```

Result:

Thus, the output was executed successfully.

Week: 11. A

SHOPPING MALL CUSTOMER SEGMENTATION

Date:

Aim:

To segment customers of a shopping mall into different groups based on their spending patterns and income using clustering.

Program Code:

```
import pandas as pd
from sklearn.cluster import KMeans
from sklearn.preprocessing import
StandardScaler # Load dataset
df =
pd.read_csv("shoppingmall_customer_segmentatio
n.csv") # Select relevant features for clustering
X = df[['Age', 'Annual_Income', 'Spending_Score', 'Gender',
'Membership_Duration']] # Convert categorical variables to
numerical
X = pd.get_dummies(X, columns=['Gender'],
drop_first=True) # Standardize the data
scaler = StandardScaler()
X_scaled =
scaler.fit_transform(X)
# Determine the optimal number of clusters using the
Elbow Method wcss = []
for i in range(1, 11):
    kmeans = KMeans(n_clusters=i, init='k-means++',
    random_state=42) kmeans.fit(X_scaled)
    wcss.append(kmeans.inertia_) # Train K-Means
clustering model
optimal_k = 5 # Choose the number of clusters based on the
elbow method kmeans = KMeans(n_clusters=optimal_k, init='k-
means++', random_state=42) df['Cluster'] =
kmeans.fit_predict(X_scaled)
# Save the clustered data
```

```
df.to_csv("shopping_mall_customers_clustered.csv",  
index=False)
```

Output:

Delimiter:

	Customer_ID	Age	Annual_Income	Spending_Score	Gender	Membership_Duration	Cluster
1	101	25	40000	65	Male	2	2
2	102	34	70000	45	Female	5	4
3	103	22	25000	80	Female	1	3
4	104	40	90000	30	Male	8	0
5	105	29	50000	55	Male	3	2
6	106	31	65000	70	Female	4	4
7	107	27	45000	60	Female	2	1
8	108	38	85000	25	Male	7	0
9	109	23	30000	85	Female	1	3
10	110	35	75000	40	Male	6	0

Result:

Thus, the output was executed successfully.

Aim:

To recommend products to customers by clustering similar customers based on their purchasing behavior.

Program Code:

```
import pandas as pd
from sklearn.cluster import KMeans
from sklearn.preprocessing import
StandardScaler # Load dataset
df =
pd.read_csv("product_recommendation.
csv") # Select appropriate features for
clustering
X = df[['product_sales', 'product_rating', 'customer_review',
'product_quality']] # Standardize the data
scaler = StandardScaler()
X_scaled =
scaler.fit_transform(X)
# Determine the optimal number of clusters using the
Elbow Method wcss = []
for i in range(1, 11):
    kmeans = KMeans(n_clusters=i, init='k-means++',
    random_state=42) kmeans.fit(X_scaled)
    wcss.append(kmeans.inertia_) # Train K-Means
clustering model
optimal_k = 5 # Choose the number of clusters based on the
elbow method kmeans = KMeans(n_clusters=optimal_k, init='k-
means++', random_state=42) df['Customer_Segment'] =
kmeans.fit_predict(X_scaled)
# Save the clustered data
df.to_csv("customer_segments.csv",
index=False)
```

Output:

Delimiter:

	product_id	product_sales	product_rating	customer_review	product_quality	Customer_Segment
1	101	500	4.5	200	85	1
2	102	300	3.8	150	70	0
3	103	700	4.9	350	90	2
4	104	250	3.5	100	60	0
5	105	600	4.7	300	88	1
6	106	400	4.2	180	75	3
7	107	800	5.0	400	95	2
8	108	200	3.2	80	50	4
9	109	550	4.6	270	83	1
10	110	450	4.0	210	78	3

Result:

Thus, the output was executed successfully.

Aim:

To analyze and segment tourists' behaviors based on their spending and travel patterns using clustering.

Program Code:

```
import pandas as pd
from sklearn.cluster import KMeans
from sklearn.preprocessing import
StandardScaler # Load dataset
df =
pd.read_csv("tourist_behaviour.csv") # Select relevant features for
clustering
X = df[['Spending_Budget', 'Trip_Frequency', 'Travel_Duration',
'Preferred_Attractions', 'Online_Bookings']]
# Standardize the
data scaler =
StandardScaler()
X_scaled = scaler.fit_transform(X)
# Determine the optimal number of clusters using the
Elbow Method wcss = []
for i in range(1, 11):
    kmeans = KMeans(n_clusters=i, init='k-means++',
    random_state=42) kmeans.fit(X_scaled)
    wcss.append(kmeans.inertia_) # Train K-Means
clustering model
optimal_k = 5 # Choose the number of clusters based on the
elbow method kmeans = KMeans(n_clusters=optimal_k, init='k-
means++', random_state=42) df['Tourist_Segment'] =
kmeans.fit_predict(X_scaled)
# Save the clustered data
df.to_csv("tourist_segments.csv",
index=False)
```


Output:

Delimiter:

	tourist_id	Spending_Budget	Trip_Frequency	Travel_Duration	Preferred_Attractions	Online_Bookings	Tourist_Segment
1	201	1500	5	7	3	1	2
2	202	2500	8	10	4	1	3
3	203	1000	3	5	2	0	4
4	204	3000	10	14	5	1	0
5	205	1800	6	8	3	1	2
6	206	1200	4	6	2	0	1
7	207	2700	9	12	4	1	0
8	208	800	2	4	1	0	4
9	209	2200	7	9	3	1	3
10	210	1400	5	7	2	0	1

Result:

Thus, the output was executed successfully.

Aim:

To segment customers into groups with similar behaviors using DBSCAN clustering.

Program Code:

```
import pandas as
pd import numpy as
np
from sklearn.preprocessing import
StandardScaler from sklearn.cluster
import KMeans
# Load the dataset
df = pd.read_csv("customer_segmentation.csv") # Replace with
your actual CSV file # Select relevant features for clustering
features = df[['Age', 'Annual_Income',
'Spending_Score']] # Standardize the data
scaler = StandardScaler()
X_scaled =
scaler.fit_transform(features) #
Apply K-Means Clustering
num_clusters = 4 # Choose optimal clusters using Elbow Method
kmeans = KMeans(n_clusters=num_clusters, random_state=42,
n_init=10) df['Cluster'] = kmeans.fit_predict(X_scaled)
# Save the segmented data
df.to_csv("customer_segmented_dataset.csv",
index=False) # Print the cluster distribution
print(df['Cluster'].value_counts())
print("Customer segmentation completed and saved as
'customer_segmented_dataset.csv'.")
```

Output:

```
2 4
0 3
1 2
3 1
```

Name: Cluster, dtype: int64

Customer segmentation completed and saved as 'customer_segmented_dataset.csv'.

Result:

Thus, the output was executed successfully.

Aim:

To analyze social network structures and group users into clusters based on their connections using DBSCAN.

Program Code:

```
import pandas as
pd import numpy as
np
from sklearn.preprocessing import
StandardScaler from sklearn.cluster
import DBSCAN
# Load the dataset
df = pd.read_csv("social_network_data.csv") # Replace
with actual file # Select relevant features
features = df[['Num_Friends', 'Mutual_Connections', 'Interaction_Frequency',
'Network_Centrality']] # Standardize the data
scaler = StandardScaler()
X_scaled =
scaler.fit_transform(features) #
Apply DBSCAN clustering
dbscan = DBSCAN(eps=0.9, min_samples=3) # Adjust parameters based on network
structure df['Cluster'] = dbscan.fit_predict(X_scaled)
# Save results
df.to_csv("social_network_clusters.csv",
index=False) # Print cluster distribution
print("Clusters found:",
df['Cluster'].value_counts()) print("Results
saved as 'social_network_clusters.csv'.")
```

Output:

Clusters found: 0 7

1 3

Name: Cluster, dtype: int64

Results saved as 'social_network_clusters.csv'.

Result:

Thus, the output was executed successfully.

Week: 13

Date:

CLASSIFY IRIS FLOWERS INTO SPECIES (SETOSA, VERSICOLOR OR VIRGINICA) BASED ON THEIR SEPAL & PETAL MEASUREMENT

Aim:

To classify Iris flowers into species (setosa, versicolor, or virginica) based on sepal and petal measurements using classification techniques.

Program Code:

```
import pandas as pd
from sklearn.model_selection import
train_test_split from sklearn.ensemble
import RandomForestClassifier from
sklearn.preprocessing import
StandardScaler
from sklearn.metrics import accuracy_score, classification_report,
confusion_matrix # Load dataset
df =
pd.read_csv("iris_flower.csv
") # Select features and
target variable
X = df[['sepal_length', 'sepal_width', 'petal_length',
'petal_width']] y = df['species']
# Standardize the
data scaler =
StandardScaler()
X_scaled =
scaler.fit_transform(X) #
Split into training and
testing sets
X_train, X_test, y_train, y_test = train_test_split(X_scaled, y, test_size=0.3,
random_state=42) # Train Random Forest model
model = RandomForestClassifier(n_estimators=100,
random_state=42) model.fit(X_train, y_train)
# Make predictions
```

```
y_pred =  
model.predict(X_test) #  
Evaluate model  
performance  
print(f"Accuracy: {accuracy_score(y_test, y_pred)}")  
print("Confusion Matrix:\n", confusion_matrix(y_test,  
y_pred)) print("Classification Report:\n",  
classification_report(y_test, y_pred)) # Display  
predictions with actual labels  
predictions_df = pd.DataFrame({"Actual": y_test, "Predicted": y_pred})
```

```
print("\nPredicted Species for Test  
Data:\n")
```

```
print(predictions_df.to_string(index  
=False))
```

Output:

Accuracy: 1.0

Confusion

Matrix:

```
[[2 0 0]
```

```
[0 3 0]
```

```
[0 0 4]]
```

Classification Report:

	precision	recall	f1-score	support
--	-----------	--------	----------	---------

setosa	1.00	1.00	1.00	2
--------	------	------	------	---

versicol or	1.00	1.00	1.00	3
----------------	------	------	------	---

virginic a	1.00	1.00	1.00	4
---------------	------	------	------	---

accuracy		1.00		9
		0		

macro avg	1.00	1.00	1.00	9
		0		

weighted avg	1.00	1.00	1.00	9
		0		

Predicted Species for Test Data:

Actual Predicted

virginica virginica

versicolor versicolor

virginica virginica

versicolor versicolor

setosa setosa

setosa setosa

virginica

virginica

virginica

virginica
versicolor
versicolor

Result:

Thus, the output was executed successfully.